

Operating Systems: Processes

CIS*3110: Operating Systems

Dr. A. Hamilton-Wright

School of Computer Science
University of Guelph

2025-01-19

Parts of a Computer



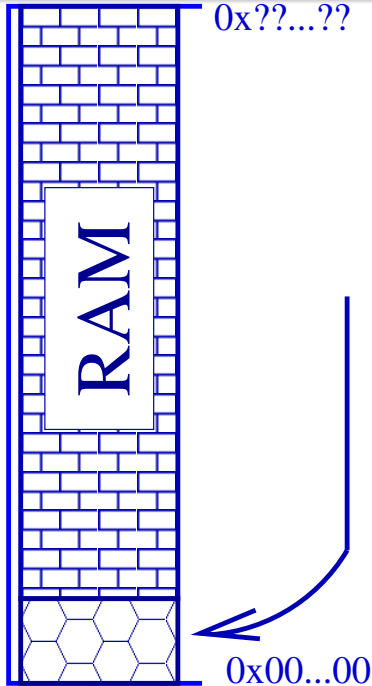
A diagram showing three components of a computer: CPU, RAM, and I/O Devices. The CPU is represented by a blue square on the left. The RAM is represented by a red square in the top right. The I/O Devices are represented by a black square in the bottom right.

CPU

RAM

I/O
Devices

Actual Physical Memory



"Zero Page"

- interrupt vector

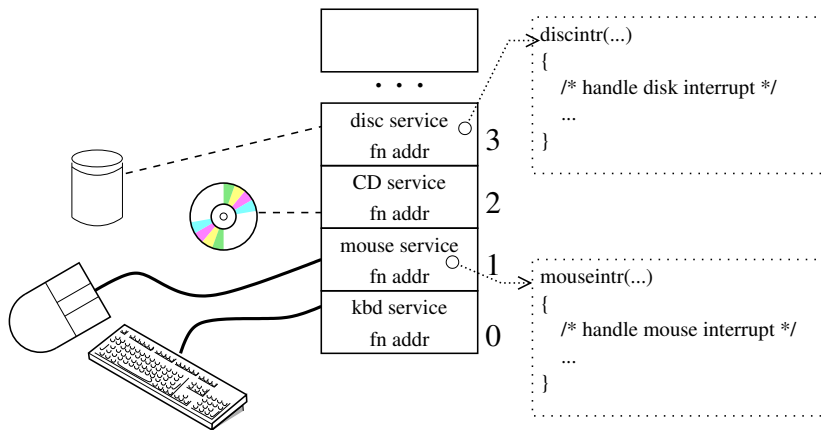
- I/O

BIOS/ROMs (if present)

Tools at the Hardware Architecture Level

- program counter (PC)
- stack pointer/frame pointer
- processor status register
- general purpose registers
- memory management unit (mapping, protection)
- kernel/user mode
- interrupts and interrupt vector

Interrupt vector links devices with service routines



- basic **instruction cycle** of the ALU of the CPU:

1. fetch instruction pointed to by the PC
 2. increment PC
 3. execute instruction
- } over and over again

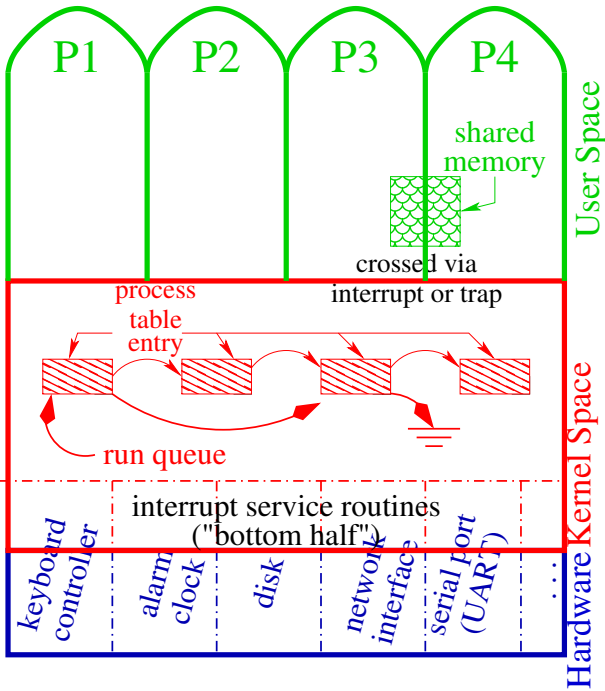
- **processor status** (usually a special register on each CPU)

user mode – can execute **most** instructions (but no I/O)

kernel mode – can execute **all** user instructions,
PLUS access special instructions and
registers that control I/O, memory access
(also called **supervisor mode**)

- **interrupt vector** (special range of locations in RAM)

- used to link hardware interrupts with ... **interrupt service routines**



Program on an Example Machine

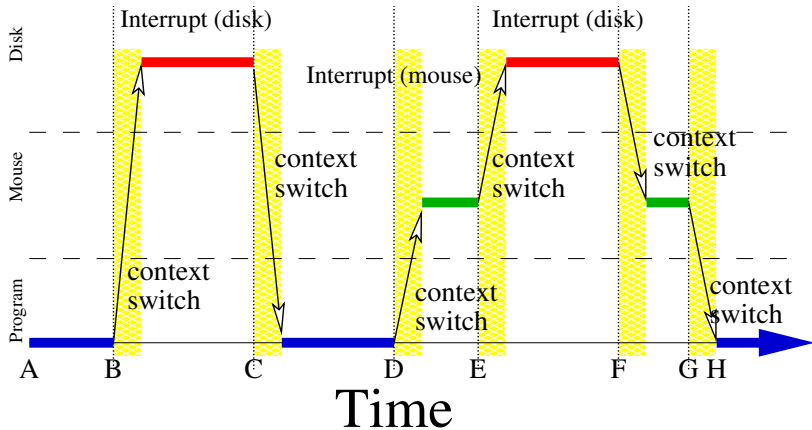
Assembler			Mach Code	Virt Addr
	load	R1, 0xA2	0x11A2	0x10
	incr	R1	0x3100	0x12
	store	R1, 0xA2	0x21A2	0x14
	jmp	xx	0x5026	0x16
yy:	decr	R2	0x4200	0x18
	
	rsb		0x7000	0x24
xx:	zero	R2	0x8200	0x26
	incr	R2	0x3200	0x28
	store	R2, 0xA0	0x22A0	0x2A
	jsr	yy	0x6018	0x2C
	incr	R2	0x3200	0x2E
	store	R2, 0xA2	0x22A0	0x30
	

Registers: R0, R1, R2, R3, SP, PC

Op Codes	
load	0x1
store	0x2
incr	0x3
decr	0x4
jmp	0x5
jsr	0x6
rsb	0x7
zero	0x8

Context Switch

- change from one running program to another “on the fly”
- need to be able to change back later
- to do this, we must save the “program state”
 - program counter
 - general purpose registers
 - status register
 - stack pointer
- each program must be resident in different parts of memory



- a hardware generated jump to subroutine (JSR) requested by an external hardware device (not the CPU)
- defined through the **interrupt vector**
 - resides somewhere in physical memory
 - is effectively an array of **function pointers**
- may be delayed (*i.e.*; disabled, possibly temporarily) by **masking off**

Interrupt Steps

Interrupt →

- save PC on stack
- switch CPU to kernel (supervisor) mode
- load PC from vector using hardware id (vector contains the start address of the interrupt service routine)
- save all registers
- continue execution

Return from Interrupt (RTI)



- restore registers
- restore saved PC value
- set CPU to previous mode

Critical Sections

- a **critical section** contains reference(s) to data shared between two or more “threads of control”
 - one may be an **interrupt**
 - this may be two or more **processes**
 - this may be two or more **threads**
- must manage the data so all parties can rely on its value
 - if we fail to do this, data may be wrong for one party, or even corrupted for all

		load	R1, #i
i++;	↔	incr	R1
		store	R1, #i

A Critical Example

```
c = 'A';  
for (i = 0; i < MAX; i++) {  
    a[i] = c++;  
    if (c > 'E')  
        c = 'B';  
}
```

Interrupt Service Routine

```
set_to_a() {  
    c = 'A';  
}
```

Modern (UNIX-type) systems: 3 types of device drivers

BLOCK devices that read/write blocks (disc, tapes, SD-cards, CD-ROM/DVD, ...)

NETWORK local area network (Ethernet, token ring), remote connection network interface (PPP, ISDN, ...)

CHARACTER everything that does not fit the above two cases (serial ports, display/keyboard, mouse, joystick, touchscreen, ...)

Device Driver Functions

Top Half: communicate with user, kernel

probe check for hardware at boot time

open called when a process does an open of a **device file** (e.g., `/dev/mouse`, `/dev/disk0`, `/dev/null` etc.)

- mostly initializes variables and maybe hardware
- generates a file descriptor in user space

close called when a process does a `close()` of the file descriptor created from `open` (usually just marks the device closed)

read called when a process does a `read(fd)` — transfers data from device to process

write called when a process does a `write(fd)` — queues data for device to output

Bottom Half: event handling

interrupt service routine

called when device generates an interrupt and does the I/O on the hardware


```

char get_keyb()
{
    while (???) {
    }

    return ???;
}

void keyb_intr()
{
    /* "magic" addresses for memory mapped I/O */
    char *keyb_reg = (char *) 0x2ff4200;
    char *keyb_status = (char *) 0x2ff4201;

    while (*keyb_status)
        ;

    ??? = *keyb_reg;
}

```

Circular Queue (Ring Buffer)

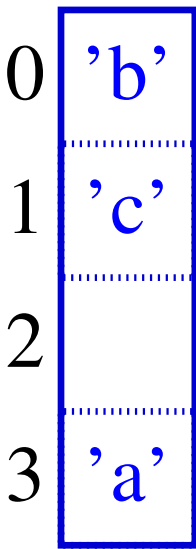
```
char Q[QSIZE];  
int numQ, startQpos, endQpos;  
numQ = startQpos = endQpos = 0;
```

Add into queue

```
Q[endQpos] = value;  
endQpos = (endQpos + 1) % QSIZE;  
numQ++;
```

Remove from queue

```
value = Q[startQpos];  
startQpos = (startQpos + 1) % QSIZE;  
numQ--;
```



Critical Sections in Device Driver

- interrupts can be masked off:

`int state = spltty()` masks interrupt(s) to prevent a second interrupt from over-writing the first

`int splx(int prevState)` returns the interrupt mask to the previous state

- process can be “put to sleep” until there is space in a buffer:

`int tsleep(void *id, int flags, char *mesg, int timeo)`
causes the current **thread** (program counter) to be **BLOCKED** until someone calls `wakeup()`

`int wakeup(void *id)` un-blocks *all* **threads** which are asleep and have the given **identifier**

Device Driver Skeleton – read

```
static char Q[QSIZE];
static int numQ = 0, startQpos = 0, endQpos = 0;
devread(char *ubuf) { /** read one character */
    int s;
    s = spltty();
    while (numQ == 0) {
        splx(s); /** can't sleep inside critical section */
        tsleep(devread, ...); /* sleep */
        s = spltty();
    }
    *ubuf = Q[ startQpos ];
    startQpos = (startQpos + 1) % QSIZE;
    numQ--;
    splx(s);
}
```


```
devintr() {  
    int s;  
    s = spltty();  
    if (QSIZE == numQ) {  
        /* queue is full ... */  
        splx(s); /* leave crit section before return */  
        return;  
    }  
    Q[ endQpos ] = GET_CHAR_FROM_DEVICE();  
    endQpos = (endQpos + 1) % QSIZE;  
    numQ++;  
    splx(s);  
    wakeup(devread);  
}
```

Running Programs

The **execution** of a **program** (an “image”) is a **process**.
The environment that the O/S provides for the execution of a program (*i.e.*; a process) is that of a **simplified** pseudo-computer.

UNIX/Mach process environment:

- **virtual** (simplified) address space
- **process context**
- no direct access to **I/O ports**



Process Table Entry (Process Context)

- the **process table** contains an entry for each current process
 - if a process has more than one **thread** there is usually one **process table entry** per thread (`ps(1)` may or may not show it this way)
- each **entry** holds all the information about a process:
 - process (thread) context
 - program counter
 - registers (general and special purpose)
 - stack pointer
 - virtual address space reference
 - scheduling information
 - **BLOCKED** -vs- **RUNNABLE** -vs- **RUNNING**
 - priority
 - time used

Process Status Listing

% ps auxw

USER	PID	%CPU	%MEM	VSZ	RSS	TTY	STAT	START	TIME	COMMAND
root	1	0.0	0.0	1368	84	?	S	Jan11	0:04	init [3]
root	2	0.0	0.0	0	0	?	SW	Jan11	0:00	[keventd]
root	4	0.0	0.0	0	0	?	SWN	Jan11	0:00	[ksoftirqd_CPU0]
root	9	0.0	0.0	0	0	?	SW	Jan11	0:00	[bdflood]
root	5	0.0	0.0	0	0	?	SW	Jan11	0:00	[kswapd]
root	1609	0.0	0.0	1368	144	?	S	Jan11	0:00	klogd -x
root	1761	0.0	0.0	3508	480	?	S	Jan11	0:00	/usr/sbin/sshd
root	1920	0.0	0.2	2320	1380	?	S	Jan11	0:00	[login]
root	1921	0.0	0.0	1352	60	tty2	S	Jan11	0:00	/sbin/mingetty
andrew	8506	0.0	0.2	4692	1448	tty1	S	09:33	0:00	-csh
andrew	9373	0.0	0.5	8196	2980	pts/7	S	10:43	0:03	vim Processes.t
andrew	9795	0.0	0.2	4680	1428	pts/8	S	11:20	0:00	-bin/csh
andrew	10154	0.0	0.1	2728	780	pts/7	R	11:36	0:00	ps auxw

Synchronization Primitives

Shared data can also be used between **processes/threads**

→ we need to manage critical sections at this level.

The two most popular **synchronization primitives** are:

mutex: bracket critical section so only 1 process allowed at a time:

```
critical_begin();  
... code in critical section ...  
critical_end();
```

semaphores: allow N -at-a-time

Co-operating sequential processes can be envisioned as a separate, dedicated processor for each – with no guarantee which runs first/faster etc.

Mutex Synchronization: Producer

```
void add_queue(char aNewChar) {  
    int oldsize;  
  
    do { /** wait until not full */  
        critical_begin();  
        oldsize = numQ;  
        critical_end();  
    } while (oldsize == QSIZE);  
  
    critical_begin();  
    Q[ endQpos ] = aNewChar;  
    endQpos = (endQpos + 1) % QSIZE;  
    numQ++;  
    critical_end();
```

Mutex Synchronization: Consumer

```
void del_queue(char aNewChar) {  
    int oldsize, tmpBuf;  
  
    do { /** wait until not empty */  
        critical_begin();  
        oldsize = numQ;  
        critical_end();  
    } while (oldsize == 0);  
  
    critical_begin();  
    tmpBuf = Q[ startQpos ];  
    startQpos = (startQpos + 1) % QSIZE;  
    numQ-;  
    critical_end();  
    return tmpBuf;  
}
```

Implementing Mutual Exclusion (in Hardware)

```
static mutexFlag = 0;

critical_begin() {
    int oldFlag;

    do {
        /** does oldflag = flag , flag = 1 */
        TEST_AND_SET(&oldFlag, &mutexFlag);
        /** as one indivisible operation */
    } while (oldFlag);
}

critical_end() {
    mutexFlag = 0;
}
```

Dekker's Algorithm (2 processes)

```
int need[2] = {0, 0};  
int turn;  
  
critical_begin(int who) {  
    need[ who ] = 1;  
    turn = !who;  
  
    while (need[ !who ] && turn != who)  
        ;  
}  
  
critical_end(int who) {  
    need[ who ] = 0;  
}
```

A **semaphore** is a synchronization primitive which has an internal integer counter supporting the following operations:

- signal** increments the internal counter

- wait** until the internal counter > 0 ...*only then* decrement the counter

Semaphore Implementation - Using Mutexes

```
SEM_WAIT(semaphore_t *s) {  
    semaphore_t oldsem;  
  
    /** wait until sem > 0 */  
    do {  
        critical_begin();  
        oldsem = *s;  
        if (*s > 0)  
            *s = *s - 1;  
        critical_end();  
    } while (oldsem == 0);  
}
```

```
SEM_SIGNAL(semaphore_t *s) {  
    critical_begin();  
    *s = *s + 1;  
    critical_end();  
}
```

- commercial/professional implementations use `tsleep()` to block the process instead of a busy-loop

Semaphore Producer/Consumer

```
available_sem = QSIZE;
used_sem = 0;
critical_sem = 1;

void add_queue(char aNewChar) {
    SEM_WAIT( &available_sem );
    SEM_WAIT( &critical_sem );
    Q[ endQpos ] = aNewChar;
    endQpos = (endQpos + 1) % QSIZE;
    SEM_SIGNAL( &critical_sem );
    SEM_SIGNAL( &used_sem );
}
```

```
void del_queue(char aNewChar) {
    int tmpBuf;
    SEM_WAIT( &used_sem );
    SEM_WAIT( &critical_sem );

    tmpBuf = Q[ startQpos ];
    startQpos = (startQpos + 1)
                % QSIZE;

    SEM_SIGNAL( &critical_sem );
    SEM_SIGNAL( &available_sem );

    return tmpBuf;
}
```


Co-Operating Sequential Processes – Examples

```
semaphore_t U = 1;    semaphore_t D = 0;    int sharedch = 'a';
```

```
up() {  
    while (1) {  
  
        printf("U-%c ", sharedch);  
        if(++sharedch > 'z')  
            sharedch = 'a';  
  
    }  
}
```

```
down() {  
    while (1) {  
  
        printf("D-%c ", sharedch);  
        if(--sharedch < 'a')  
            sharedch = 'a';  
  
    }  
}
```

Co-Operating Sequential Processes – Example #1

```
semaphore_t U = 1;    semaphore_t D = 0;    int sharedch = 'a';
```

```
up() {  
    while (1) {  
        SEM_WAIT( &U );  
  
        printf("U-%c ", sharedch);  
        if(++sharedch > 'z')  
            sharedch = 'a';  
        SEM_SIGNAL( &D );  
    }  
}
```

```
down() {  
    while (1) {  
        SEM_WAIT( &D );  
  
        printf("D-%c ", sharedch);  
        if(--sharedch < 'a')  
            sharedch = 'a';  
        SEM_SIGNAL( &U );  
    }  
}
```

Output: U-a D-b U-a D-b U-a ...

Co-Operating Sequential Processes – Example #2

```
semaphore_t U = 2;    semaphore_t D = 0;    int sharedch = 'a';
```

```
up() {  
    while (1) {  
        SEM_WAIT( &U );  
  
        printf("U-%c ", sharedch);  
        if(++sharedch > 'z')  
            sharedch = 'a';  
        SEM_SIGNAL( &D );  
    }  
}
```

```
down() {  
    while (1) {  
        SEM_WAIT( &D );  
        SEM_WAIT( &D );  
        printf("D-%c ", sharedch);  
        if(--sharedch < 'a')  
            sharedch = 'a';  
        SEM_SIGNAL( &U );  
        SEM_SIGNAL( &U );  
    }  
}
```

Output: U-a U-b D-c U-c U-d D-e U-d U-e D-f U-e U-f D-g U-f ...

Co-Operating Sequential Processes – Example #3

```
semaphore_t U = 0;    semaphore_t D = 2;    int sharedch = 'a';
```

```
up() {  
    while (1) {  
        SEM_WAIT( &U );  
        printf("U-%c ", sharedch);  
        if(++sharedch > 'z')  
            sharedch = 'a';  
        SEM_SIGNAL( &D );  
    }  
}
```

```
down() {  
    while (1) {  
        SEM_WAIT( &D );  
        SEM_WAIT( &D );  
        printf("D-%c ", sharedch);  
        if(--sharedch < 'a')  
            sharedch = 'a';  
        SEM_SIGNAL( &U );  
        SEM_SIGNAL( &U );  
    }  
}
```

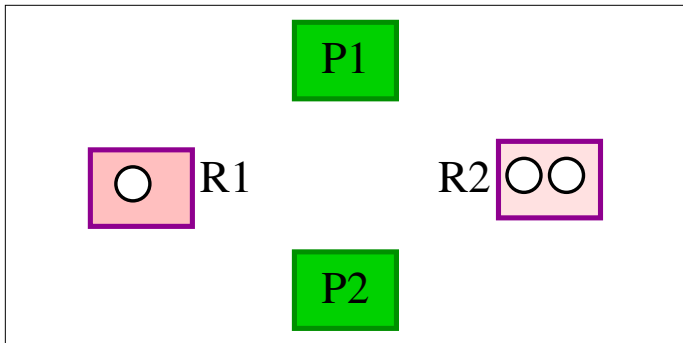
Output: D-a U-a U-b D-c U-c U-d D-e U-d U-e D-f U-e U-f D-g ...

Semaphore Hints:

- all semaphores must be properly initialized
- always work from the definition of `wait_sem()` and `signal_sem()`
- normally, there will be a `wait_sem` with a corresponding `signal_sem()` on the same semaphore
- the order of `wait_sem()` operations is quite important; the order of **signal** less so

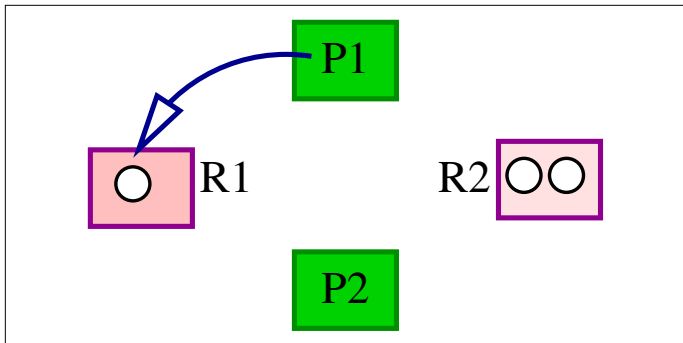
Deadlock

- deadlock occurs when processes *hold* resources and they *require* resources that other processes already hold
- shown by the presence of a cycle in the request allocation graph



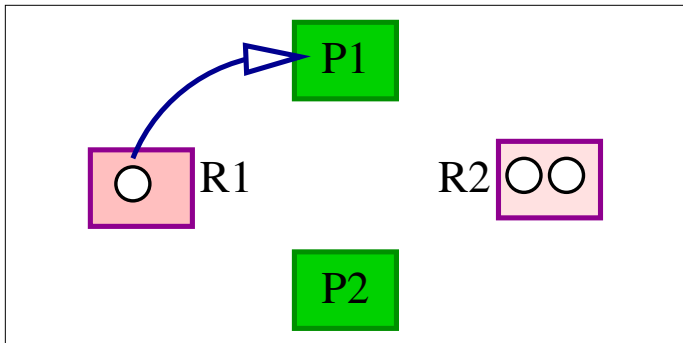
Deadlock

- deadlock occurs when processes *hold* resources and they *require* resources that other processes already hold
- shown by the presence of a cycle in the request allocation graph



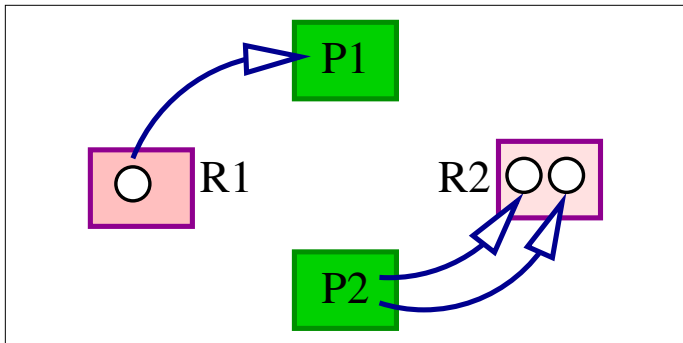
Deadlock

- deadlock occurs when processes *hold* resources and they *require* resources that other processes already hold
- shown by the presence of a cycle in the request allocation graph




Deadlock

- deadlock occurs when processes *hold* resources and they *require* resources that other processes already hold
- shown by the presence of a cycle in the request allocation graph




Deadlock

- deadlock occurs when processes *hold* resources and they *require* resources that other processes already hold
- shown by the presence of a cycle in the request allocation graph



Deadlock

- deadlock occurs when processes *hold* resources and they *require* resources that other processes already hold
- shown by the presence of a cycle in the request allocation graph



Java `synchronized` keyword

- The `synchronized` keyword can be used on a whole method or a block to provide a critical section
- method locking:

```
public class SyncCount {  
    private int count = 0;  
    public synchronized void increment() { count++; };  
    public synchronized int getValue() { return count; };  
}
```

- block locking:

```
public void someMethod() {  
    while(keepGoing()) {  
        synchronized(someObject) {  
            count++;  
        }  
    }  
}
```

Shared Memory blocks of shared process image

Semaphores/Mutexes integer data values for flow control

Signals a software “interrupt”

Pipes an imitation file shared by two processes

- used to indicate that some important event has happened
- sequence of events
 - process A is running and performing some task
 - process B sends a signal to process A
 - process A interrupts whatever it is doing ...
 - call a signal handler (a function) ...
 - returns to whatever it was doing before
- there is no direct notification to process B that anything has been done

Signal Values

Signals are simply integer valued actions, similar to interrupts.
Some common signals are:

1	HUP	hangup	BUS	bus error
2	INT	interrupt	SEGV	segmentation violation
3	QUIT	quit	PWR	power failure
4	ILL	illegal instruction	TERM	terminate process
5	TRAP	trace trap	USR1	user defined #1
9	KILL	kill process	USR2	user defined #2

Signal Example

```
#include <signal.h>

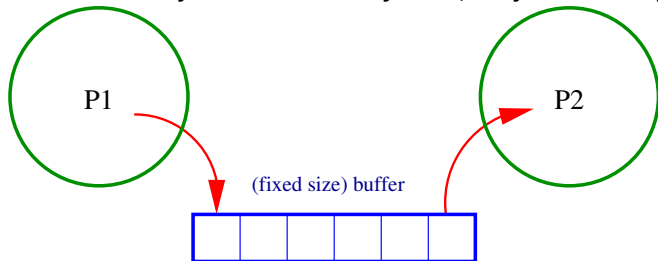
void signalHandler(int code)
{
    fprintf(stderr,
              "Interrupt caught — cleaning up\n");
    performCleanup();
    exit(1);
}

void main()
{
    signal(SIGINT, signalHandler);

    :
```


Pipes

A **pipe** is a file descriptor pair shared between processes, attached to a buffer common to both processes. The buffer is not stored anywhere in a filesystem, only in memory.



Basic process scheduler loop:

- when a process can no longer run (exits, waiting for I/O or time-slice used up)
 - mark process RUNNABLE or blocked (and why – blocked status)
 - save process context in process table entry
- search process table for another RUNNABLE process
 - mark this process running
 - restore process context for this process
- the system now runs this process, picking up where it left off (just as with a hardware interrupt)

Scheduling Algorithm or “Queuing Discipline”

FIFO

- if jobs/processes/...repeatedly join the queue this is typically called “round robin.”
- typically, round robin is preemptive

SJF

- shortest job first
- allow the job with the shortest running time to go next (jump to the head of the queue)
- not preemptive

A preemptive discipline permits a job to be preempted before it completes

Scheduling Algorithms/Queueing Disciplines

- FIFO and round robin are generally considered to be **fair**, as jobs cannot be stuck waiting forever
- SJF is optimal w.r.t. throughput:
 - *iff* you always run the job that will finish in the shortest time.
 - will always finish more jobs in a given elapsed time than any other ordering
 - it is **unfair**, and impractical for processes (can be approximated)

Scheduling Algorithms/Queuing Disciplines

In general:

- preemptive schedulers will perform better (*i.e.*; smoother response) than non-preemptive schedulers and are **required** for interactive systems
- schedulers will normally bias towards short jobs, such as updating an editor after a keystroke
- operators/users will be able to set/adjust priorities for jobs, see `nice(1)`
 - dynamic priority is updated during context switch based on last run time
- there will always be a trade-off between coarse -vs- fine granularity of the time-slices:
 - coarse** : low overhead/poor interactive response
 - fine** : higher overhead/better interactive response

Scheduling Algorithms/Queueing Disciplines

- if preemption occurs extremely frequently such that each job gets an infinitesimally small time each time it runs, the discipline is called “processor sharing”
- whereas, if the preemption time is infinitely large, the discipline is called FIFO

Scheduling Algorithms/Types of Tasks


batch insensitive to gaps in runtime; get overall task done

interactive sensitive to runtime gaps; typically lots of I/O

real-time critically sensitive to runtime gaps

- demand fixed fraction of CPU/unit time

Priority Queues (VMS/Windows)



Some systems vary the priority dynamically, others do not