

Operating Systems: I/O

CIS*3110: Operating Systems


Dr. A. Hamilton-Wright

School of Computer Science
University of Guelph

2024-12-01


I/O, I/O it's off to work we go . . .

I/O Subsystem:





- DMA → block device \rightsquigarrow block device driver
- data transfer using data bus → character device

I/O Subsystem:



Disk



A disk address is three-dimensional (θ, r, z); data is stored:

in sectors θ
within cylinders r
on tracks (surfaces, heads) z

- **seeking** selects a cylinder
- **rotation** provides access to all sectors

Disk Parameters

Example:

Seagate Cheetah 10K.6
(SCSI bus)

10,000 RPM

→ 6 ms/revolution

→ 3 ms average
latency

4 physical platters (discs)

8 heads

49,854 cylinders

772 sectors/track
(@ 512 bytes/sector)

5 ms avg seek time

$$8 \times 49,854 \times 772$$

$$= 307,898,304 \text{ sectors @ } 1/2 \text{ k each}$$

$$= 153,949,152 \text{ k} \approx 146.8 \text{ Gb}$$

Data Rate:

$$772 \times 512 = 395,264 \text{ bytes/track}$$

$$\therefore 395,264 / 6 \text{ ms}$$

$$= 65,877 \text{ bytes per ms}$$

“disk-to-buffer” speed

“buffer-to-computer” speed that of
controller/memory bus – currently

~2.5Gbit/s (SCSI-Ultra320)

~3Gbit/s (SATA)

$$772 \times 8 = 6,176 \text{ sectors/cylinder}$$

Minimize Seek Distance

Q: Which is better?

FIFO (request order)


SJF do the request which is in the closest cylinder

SCAN (elevator algorithm) shortest seek distance in the direction you are travelling

Q: What size should the blocks be?

- larger blocks – more waste at EOF, larger payoff when block is found
- smaller blocks – less EOF waste, but likely more indexing data

FAT (File Allocation Table) based filesystems




FAT File System

Block Size: 1024

Data blocks start
after root dir

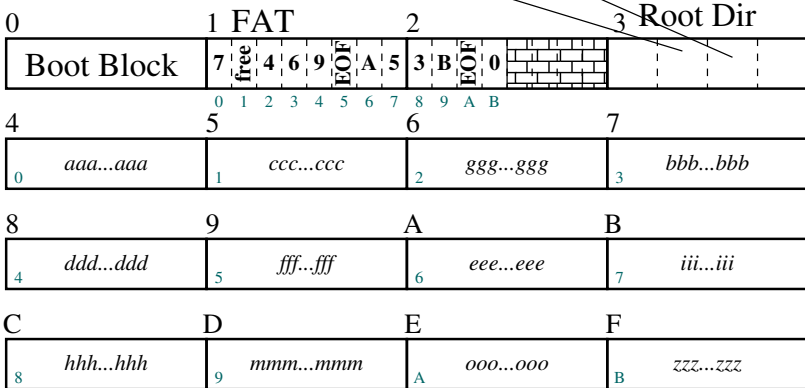

| | | | |
|--|---|--|--|
| | Name ABC.DAT Size 7200 Modify Time First Block 2 | Name A2.C Size 5000 Modify Time First Block 8 | |
|--|---|--|--|



FAT File System

Block Size: 1024

Data blocks start
after root dir




FAT File System

Block Size: 1024

Data blocks start
after root dir

| | | |
|---|--|--|
| Name ABC.DAT Size 7200 Modify Time First Block 2 | Name A2.C Size 5000 Modify Time First Block 8 | |
|---|--|--|




FAT File System

Block Size: 1024

Data blocks start after root dir

| | | |
|---|--|--|
| Name ABC.DAT Size 7200 Modify Time First Block 2 | Name A2.C Size 5000 Modify Time First Block 8 | |
|---|--|--|




FAT File System

Block Size: 1024

Data blocks start after root dir

| | | |
|---|--|--|
| Name ABC.DAT Size 7200 Modify Time First Block 2 | Name A2.C Size 5000 Modify Time First Block 8 | |
|---|--|--|




FAT File System

Block Size: 1024

Data blocks start
after root dir

| | | |
|---|--|--|
| Name ABC.DAT Size 7200 Modify Time First Block 2 | Name A2.C Size 5000 Modify Time First Block 8 | |
|---|--|--|



Data blocks start
after root dir



MBR:

| | |
|----------------------|-----------------------|
| (boot) jump location | OEM name |
| sector size | sectors per block |
| number of FATs | # of root dir entries |
| total num sectors | media type |
| sectors per track | num heads |

Directory Entry

| |
|------------------------|
| file name [8] |
| file extension [3] |
| attributes |
| modification date/time |
| file length (bytes) |
| first block |

FAT : Floppy Disk Cluster/Block Sizes

| <i>Drive Size</i> <i>(logical volume)</i> | <i>FAT Type</i> | <i>Sectors</i> <i>Per Cluster</i> | <i>Cluster</i> <i>Size</i> |
|--|-----------------|--------------------------------------|-------------------------------|
| 360K | 12-bit | 2 | 1K |
| 720 | 12-bit | 2 | 1K |
| 1.2Mb | 12-bit | 1 | 512b |
| 1.44Mb | 12-bit | 1 | 512b |
| 2.88Mb | 12-bit | 2 | 1K |

From Microsoft KB article 314878,
<http://support.microsoft.com/kb/314878>

FAT : Hard Drive Cluster/Block Sizes

| <i>Drive Size (logical volume)</i> | <i>FAT Type</i> | <i>Sectors Per Cluster</i> | <i>Cluster Size</i> |
|--|-----------------|--------------------------------|-------------------------|
| 0-15Mb | 12-bit | 8 | 4K |
| 16-127Mb | 16-bit | 4 | 2K |
| 128-255Mb | 16-bit | 8 | 4K |
| 256-511Mb | 16-bit | 16 | 8K |
| 512-1023Mb | 16-bit | 32 | 16K |
| 1024-2047Mb | 16-bit | 64 | 32K |
| 2048-4095Mb | 16-bit | 128 | 64K |

FAT : “Large” Hard Drive Cluster/Block Sizes

| <i>Drive Size</i> <i>(logical volume)</i> | <i>FAT Type</i> | <i>Sectors</i> <i>Per Cluster</i> | <i>Cluster</i> <i>Size</i> |
|--|-----------------|--------------------------------------|-------------------------------|
| 4–8Gb | 16-bit | 256 | 128K |
| 8–16Gb | 16-bit | 512 | 256K |

- not available prior to WinNT4.0, as reflected in the KB article (they are actually available afterwards)
- this is as far as basic FAT will go:
 $2^{16} \times (512 \times 512) = 17179869184 = 16384 \times (1024 \times 1024)$

FAT32 Cluster Sizes

| <i>Drive Size (logical volume)</i> | <i>Cluster Size</i> |
|--|-------------------------|
| 16-64Mb | 512b |
| 64-128Mb | 1kb |
| 128-256Mb | 2kb |
| 256Mb-8Gb | 4kb |
| 8Gb-16Gb | 8kb |

exFAT Cluster Sizes

| <i>Drive Size (logical volume)</i> | <i>Cluster Size</i> |
|--|-------------------------|
| 7-256Mb | 4kb |
| 256Mb-32Gb | 32kb |
| 32Gb-256Tb | 128kb |

NTFS : Cluster/Block Sizes (WinNT4.0 and up)


| <i>Drive Size (logical volume)</i> | <i>Sectors Per Cluster</i> | <i>Cluster Size</i> | <i>Max Offset</i> |
|--|--------------------------------|-------------------------|--|
| 7Mb-16Tb | 8 | 4k | $2^{32} \times 8 \times 512\text{b}$ $= 2^{44}\text{b} = 16\text{Tb}$ |
| 16Tb-32Tb | 16 | 8k | $2^{32} \times 16 \times 512\text{b}$ $= 2^{45}\text{b} = 32\text{Tb}$ |
| 32Tb-64Tb | 32 | 16k | |
| 64Tb-128Tb | 64 | 32k | |
| 128Tb-256Tb | 128 | 64k | $2^{32} \times 128 \times 512\text{b}$ $= 2^{48}\text{b} = 32\text{Tb}$ |

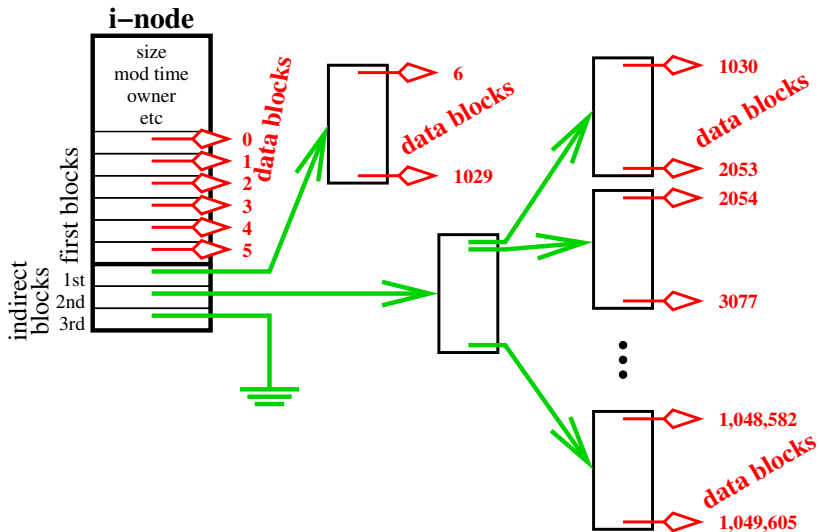
From Microsoft KB article 140365,
<http://support.microsoft.com/kb/140365>

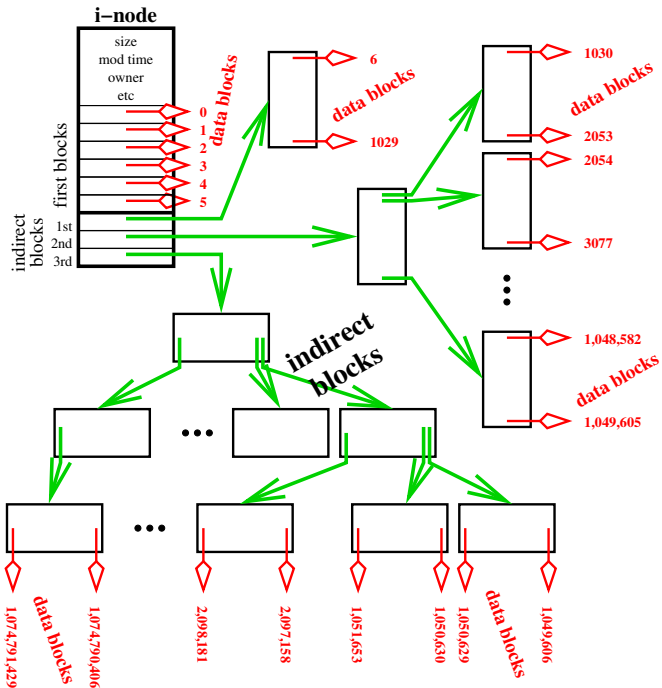
- file blocks are referenced in a **broad tree** whose root is an **index node** (also called an “**i-node**”)
- directories are stored in **files**
- the **root directory** is the **root** of the **filesystem tree**; the i-node of the root directory is in the **superblock**

- Allocated to the first block of the filesystem and repeating on a fixed pattern
- Fields:
 - offset of first data block
 - file system size
 - i-node for FS root
 - head of free i-node list
 - head of free block list
 - device info (C/H/S etc.)
 - dirty flag
 - blocks/cylinder
 - volume label
 - num free i-nodes
 - num free blocks


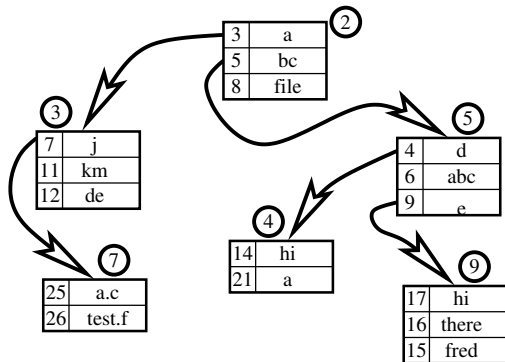
| Type | Field name |
|---------|--------------------------------------|
| long | file mode |
| long | link count |
| long | owner id |
| long | group id |
| di_size | file size |
| time_t | last accessed |
| time_t | last modified |
| time_t | i-node accessed |
| | <i>array of file block addresses</i> |








UFS Example



i – number of i-node referencing the data shown with 'dir' mode bit set

The UFS Freelist




Unix Filesystem Performance

McKusick, Marshall K. *et al*, "A Fast File System for UNIX", *Computer Systems*, 2(3):181-197, 1984.

| <i>Space Used</i> | <i>Waste</i> | <i>Organization</i> |
|-------------------|---------------|--|
| 775.2 Mb | 0.0 % | Data only, no separation between files |
| 807.8 Mb | 4.2 % | Data only, each file starts on 512 byte boundary |
| 828.7 Mb | 6.9 % | Data + inodes, 512 byte block old UFS |
| 866.5 Mb | 11.8 % | Data + inodes, 1024 byte block old UFS |
| 948.5 Mb | 22.4 % | Data + inodes, 2048 byte block old UFS |
| 1128.3 Mb | 45.6 % | Data + inodes, 4096 byte block old UFS |

Speeding up a File System: Increase block size

- but this increases wasted space!
- so collect “fragments” into a single block




Unix Filesystem Performance, cont'd

| <i>Throughput</i> | | | <i>kbytes/sec</i> |
|-------------------|-----|------|-------------------|
| read | Old | 1024 | 29 |
| | New | 4096 | 221 |
| | New | 8192 | 233 |
| write | Old | 1024 | 48 |
| | New | 4096 | 142 |
| | New | 8192 | 215 |

Speeding up a File System: Reduce seek distance

- allocate blocks close together when possible
- one way is to use **extents** — (attempt to) keep file in contiguous blocks
 - DOS/Windows “defragging” operation laboriously reorders files to achieve this
 - operation is expensive → user initiated
 - degradation over time
- another way is to use **cylinder groups** and allocate all blocks for a file (both data and metadata) from the same cylinder group



Allocation to Cylinder Groups

- growing an i-node
 - allocate within cylinder group for a while (ideally contiguously), *then*
 - spill over into other cylinder groups every few megabytes
- allocate a new i-node (regular file)
 - in same cylinder group as the directory it is in
- allocate a new i-node (directory)
 - in a different cylinder group with lots of free i-nodes and few directories

Reliability -vs-Performance

Reliability

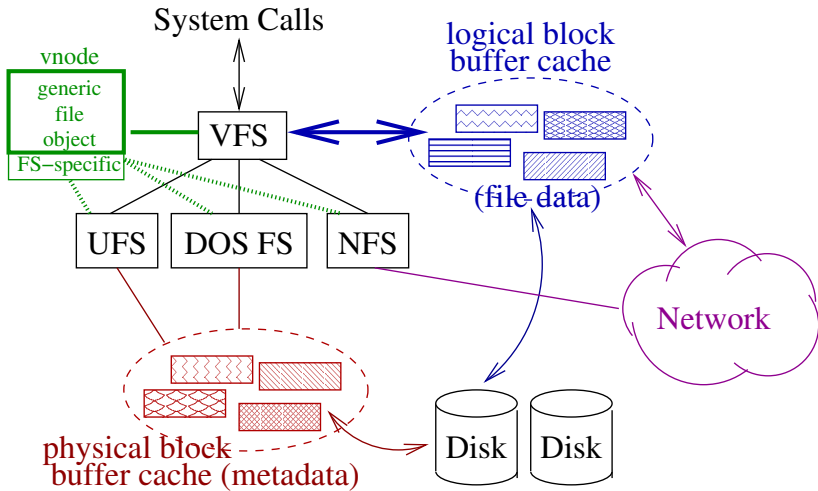
- redundant information
- synchronous writing

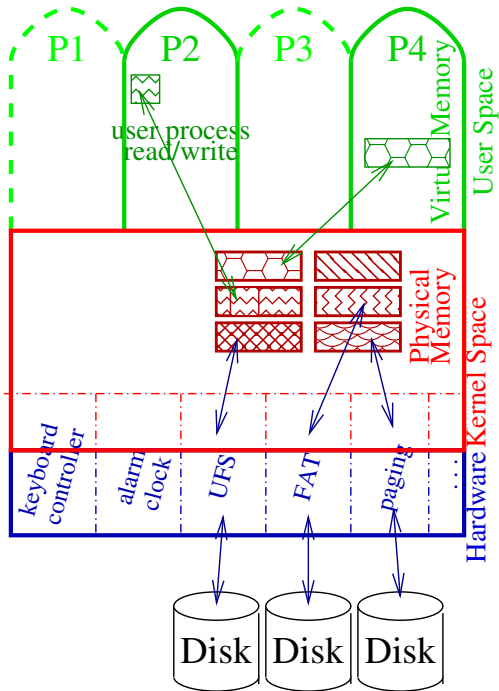
Performance

- delayed writing
- increased risk of structural damage when a crash occurs

BSD file system:

- | | |
|---|--|
| <ul style="list-style-type: none">• write new i-node before the directory entry that points to it• remove directory entry before deallocating i-node• write deallocated i-node to disk before freeing data blocks | <ul style="list-style-type: none">• also, speed up writing by<ul style="list-style-type: none">• collating several writes into one• keep the writes close together• do the writes when the disk isn't too busy |
|---|--|






- calculate block # (divide byte offset/block size)
- look in logical buffer cache for the block #
- if found
 - copy data to user process data area
- else
 - *next ...*

(else)

- allocate a block in the logical buffer cache for the block #
- perform read of the block for the file system
 - use meta-data for the file system to map file block # \Rightarrow disk block #
 - if valid data not on the disk \rightarrow fill block with zeros
 - issue disk block read to the disk controller hardware via device driver
 - wait for I/O operation to complete (*i.e.*; wait for interrupt)
 - copy data to user process data area

- calculate block # (divide byte offset/block size)
- look in logical buffer cache for the block #
- if *not* found:
 - allocate a buffer for the block #
 - if valid data not on the disk → fill block with zeros
 - else if the write won't fill the entire buffer
 - perform read of the block from the file system (same as read)
- *next ...*

- copy data from user process data area to buffer



- mark the buffer modified
- schedule a write of the block #
 - use metadata to map file block # to disk block #
 - issue disk block write to the hardware disk device controller via the device driver

Writing Schedules

Scheduling can be:

synchronous → waits for write to complete on disk

asynchronous → start write, but don't wait for completion

“write back” → do write later (*a.k.a.* “delayed write”)

Read Ahead:


- try and guess the next block that the process might read, and read it into the buffer cache in anticipation

Write Behind:

- start now, but don't wait for completion (asynchronous)

OR

- do it later (delayed write)
 - could result in several writes being combined into one



- metadata might be repeatedly modified (file length/modification time updates in particular)

SSD : Solid State Drives

- finite-write technology based on a erase-to-write technology:
 - MLC: “multi-level cell” $\propto 10^3$ total writes before failure
 - SLC: “single-level cell” $\propto 10^4$
- wear-levelling: “it is better to spread writes across drive”
 - limited cycles per erase block
 - erase block size?
 - typical filesystem (re)writes meta-data many times
 - SSD moves data around \rightarrow fragmentation
- actual solutions:
 - write tricks
 - change FS algorithms?
 - change reported data (access times)?

- if a power fails or system crashes without unmount, recovery can be quite costly
- “journaling” takes a database (transaction) approach to filesystem updates
- each write action is a **transaction**, recorded in **journal**
- each transaction is not complete until filesystem integrity is maintained (i-lists, dir entries + blocks all updated)

Journaling Protocol

- record writes to a circular queue, schedule a write
- if queue fills, block
- one physical write at a time, from head of queue
- when write completes (data + metadata) remove write from queue (journal)

Result: actions of write are reproducible

during recovery, knowledge of tasks being performed is available to reconstruct actions on data + metadata.

recovery consists of completing all actions listed in **journal**

Different schemes, differ in **currency**:

regions - ResiserFS

- transactions are recorded in offset+length structures
- may span disk blocks
- may not start/end on block bounds

blocks - ext3

- transactions are recorded in terms of blocks
- more efficient
 - **write-behind**
 - block alignment factors
 - code is simpler as block based structures exist in FS code already

Locks : (File) Data Access Synchronization

- (file) lock → **region** of a file (offset + extent)
- raises issue of intersection, inclusion
- two types (four names):
 - Read/Shared Lock :
 - multiple locks can exist for same region
 - prevents exclusive lock from being set
 - Write/Exclusive Lock :
 - complete access (to region)
- when multiple lock types are in wait and lock comes free,
who gets it?

Partitions


- divide a large disk into several smaller logical volumes of non-overlapping **extents**
- based on a table in cylinder 0 that indicates the begin/end location of each partition, and usually a filesystem type flag
- some (older) schemes partition in terms of cylinders, others can take any sector numbers

Partitions – DOS style

- DOS style partitions are set up inside the MBR
 - “primary” partition table has four slots
 - partition has an extent and a type: “primary” or “extended”
 - an “extended” partition is a container for “logical” partitions stored outside the MBR
 - (some tools work badly if there is more than one extended partition, and some boot manager cannot find filesystems in logical partitions)

- UFS partitions are set up in the first sector of the device, and typically have 8 entries:
 - a: the root filesystem;
 - b: the “swap” partition (paging area);
 - c: span - raw access to entire disk
 - d-h: general use (typically /usr, /var, /home and others)

RAID 0: “striping” or “byte/block splitting”




block-level “striping” → split the data into N parts, each write is $1/N$ original size

advantage: speed

cost: failure of any disk is fatal

overhead: none

RAID 1: “mirroring”



disk-level mirroring → keep a spare copy

advantage: reliability (single-disk failure)

cost: reduces total speed

overhead: double the space requirements

RAID 2,3,4,5: parity plans

- parity calculation allows us to catch 1-bit errors (RAID 2 - never used)
- drives self-detect failure, so we save one bit on RAID 2 plan (RAID 3)
- we can store all the parity bits for a block together (RAID 4)
- if we distribute all this across the drives, no single device is involved in every write (RAID 5)

RAID 5: block-interleaved distributed parity



parity blocks distributed across other disks

advantage: reliability (single-disk failure)

cost: two disks involved in all I/O

overhead: 1 parity per N data + calculation

parity for block m stored on disk $(m \bmod N) + 1$

RAID 6: P+Q redundancy




like level 5 but more parity \rightarrow tolerate more failures

advantage: reliability (multi-disk failure)


cost: failure level + 1 disks involved in all I/O

overhead: variable \propto redundancy (usually) based on Reed-Solomon codes (data values form coefficients for a polynomial; store values + samples to allow reconstruction)

(The general plan is to use data values as coefficients for a polynomial, and store some points on the resulting line; the line may be interpolated for reconstruction of original coefficients)



RAID 0+1: speed + safety = \$\$\$



do *both* RAID 0 and RAID 1 simultaneously

advantage: speed + reliability (single-disk failure)

cost: best case \mapsto 0 provides performance, 1 provides reliability

overhead: double the disk space needed